

Beyond Sanitizers: LLM-Guided Refinement for Differential JSON Deserialization Testing in Dart

Anonymous Author(s)
Submission under double-blind review

Abstract—A prior black-box, LLM-guided grammar-fuzzing pipeline raised a native C parser’s acceptance rate from 58.2% to 97.1% using only coverage-free feedback, and named a memory-safe target – where “crash” has no sanitizer-based meaning – as future work. We take up that challenge, re-targeting the identical refinement loop at four widely-used Dart/Flutter JSON deserialization paths (hand-written parsing, `json_serializable`, `freezed`, and `built_value`) and replacing the crash oracle with two that need no crash at all: cross-implementation differential divergence, and a ground-truth numeric round-trip check. In seeded $n=5$ -per-arm comparisons, refinement reproduces the original result on a like-for-like acceptance objective (50.0% to a 92.0% mean, exact $p \approx 0.0079$), but on the differential objective it finds divergence in only 2.9% of generated documents against 22.2% for a static, schema-aware generator. Lifting the sandbox’s JSON-library restriction, richer feedback, incremental mutation, and a larger proposer model all fail to close this gap. Giving the proposer four observations from the thirteen-case manual characterization the static generator was designed from closes it entirely: in a pre-registered, replicated comparison, refinement then reaches a 28.8% mean ($p \approx 0.013$ against score-only), matching the static generator – yet it never finds a divergence beyond what it was told. The gap is domain knowledge, not search. The campaigns also expose two practitioner-actionable defects: `json_serializable` and `freezed` silently saturate out-of-range integers to int64 bounds (5.7% of 43,336 checks), and `built_value` silently defaults a missing required list to empty.

Index Terms—fuzzing, differential testing, property-based testing, large language models, JSON deserialization, Dart, Flutter, black-box testing

I. Introduction

A black-box, LLM-guided grammar-fuzzing pipeline can be evaluated along two different generalization axes. One is swapping the target implementation while keeping everything else fixed – a prior study did this, reusing an identical pipeline unmodified against a second, independently implemented C JSON parser, and found the mechanics of the loop (harness contract, proxy signal, sandbox) transferred with no redesign [1]. The other axis, which that same study named explicitly as unaddressed future work, is swapping the oracle itself: every component of that pipeline’s failure signal – a sanitizer-instrumented build, a subprocess’s exit code, a fatal signal – assumes a native process that can crash in the memory-unsafe sense. A memory-safe language has no such crash to look for, and “failure” has to be redefined around something else entirely: uncaught exceptions, silently violated API

contracts, or values that decode incorrectly without any exception at all.

This paper takes up that second axis against a target chosen for practical relevance: JSON deserialization in Dart, the language Flutter mobile and web applications are written in. Porting the crash oracle onto a single, well-audited Dart decoder would likely reproduce the uninformative negative result the original study already reported for its second C target, so we design two oracles that do not depend on a crash at all. The first is differential divergence: every generated document runs through four independently-implemented, widely-used Dart JSON deserialization paths – hand-written parsing, and the `json_serializable`, `freezed`, and `built_value` codegen frameworks – and disagreement between them is the signal. The second is a ground-truth numeric round-trip check: JSON permits arbitrary-precision integer literals, and we check whether a decoded value preserves the exact magnitude of the literal text the generator wrote.

Retargeting the oracle changes what the refinement loop is asked to optimize, and we treat that as an empirical question. RQ1: reusing the original acceptance-rate proxy and loop essentially unmodified, does LLM-guided refinement still raise a Dart JSON generator’s rate of producing syntactically valid documents? RQ2: does malformed or boundary-case input surface divergence across the four paths, and does refinement – with its prompt corrected to make divergence the explicit objective – find more of it than a static, schema-aware generator? RQ3: does numeric decoding preserve exact integer values at magnitudes JSON permits but native integer types do not, measured at scale?

We answer RQ1 affirmatively and RQ2 negatively – the same mechanism, proposer model, and iteration budget transfer cleanly to one objective and not to the other – and then trace the RQ2 gap to its cause: not the loop’s ability to search, but the domain knowledge the static generator’s designer had and the loop did not. This paper contributes:

- Two coverage-free, crash-free oracles for a memory-safe target – cross-implementation differential divergence and ground-truth numeric round-tripping – and a two-stage harness (a shared JSON-syntax gate, then four independent schema-level decodes) that computes them automatically.
- A seeded $n=5$ -per-arm replication (RQ1) showing

the original refinement mechanism reproduces its acceptance-rate effect on a new runtime and language (50.0% to a 92.0% mean, exact two-sided $p \approx 0.0079$).

- A comparison (RQ2) showing the same mechanism underperforms a static generator by an order of magnitude on the divergence objective (2.9% vs. 22.2% mean, $p \approx 0.0079$), and three follow-ups that test rather than assume the reason: removing generation overhead (import json allowed) leaves the rate unchanged, richer category feedback makes it significantly worse ($p \approx 0.032$), and mutate-not-rewrite refinement does not replicate a promising held-out seed – localizing the gap to targeting.
- A pre-registered, replicated knowledge-parity experiment showing the gap is domain knowledge: given four observations from a thirteen-case manual characterization, refinement matches the static generator (28.8% mean, $p \approx 0.013$ against score-only), while a larger proposer without them does not improve (gpt-4.1, 1.91%, $p \approx 0.31$) – and the loop never finds divergence beyond what it was told.
- A methodological note: an initial prompt found zero divergence across three full seeded runs while driving up unrelated rejection counts; a principled revision, validated on a held-out seed before further spend, recovered a consistent effect. We report the failed design alongside the fix.
- Two practitioner-actionable defects found with no LLM involved and confirmed at scale (RQ3): `json_serializable/freeze` silently saturate out-of-range integers to int64 bounds (5.7% mismatch across 43,336 checks), and `built_value` silently defaults a missing required list to empty – plus a released pipeline, harness, and full campaign artifact set.

II. Related Work

a) Coverage-guided fuzzing.: American Fuzzy Lop [2] and its successor AFL++ [3] popularized mutation-based, coverage-guided fuzzing: a target is instrumented to report which code ran, and the fuzzer retains and mutates inputs that reach new coverage. Manès et al.’s survey [4] gives a taxonomy of this and related approaches. Coverage-guided fuzzers are highly effective when instrumentation is available, but instrumentation itself is exactly what a memory-safe, JIT/AOT-compiled runtime such as Dart’s makes both less necessary (there is no memory-safety class of bug for a sanitizer to catch) and less available in the form these tools assume; none of this family targets the failure modes this paper is concerned with.

b) Grammar-based and generation-based fuzzing.: Generating inputs directly from a grammar, rather than mutating bytes, is a much older idea: Purdom’s sentence generator [5] produces short strings that exercise every grammar production, originally for testing and debugging parsers themselves. More recent tools combine a grammar with coverage feedback: Superion [6] extends a greybox

fuzzer with grammar-aware mutation operators so mutated inputs remain syntactically valid, and Nautilus [7] similarly pairs a grammar-based generator with coverage-guided mutation to reach deeper parser states. Both still assume coverage instrumentation is available; the pipeline this paper extends removes that assumption entirely.

c) Property-based testing.: QuickCheck [8] introduced property-based testing: a developer writes generators for values of a given shape and a property that should hold for all of them, and the tool searches for a counterexample. Hypothesis [9], which the pipeline this paper reuses builds directly on, is a Python implementation of the same idea, with composable strategies expressive enough to encode both a formal grammar (our RQ1 generator) and an adversarial perturbation catalog over a target-specific schema (our RQ2 generator, Section III).

d) LLMs for fuzzing and test generation.: LLMs have recently been used to generate or mutate fuzzing inputs directly. TitanFuzz [10] uses an LLM to generate and mutate programs for fuzzing deep-learning library APIs, relying on the LLM’s prior exposure to real usage examples of those APIs rather than a formal grammar. Fuzz4All [11] generalizes this to arbitrary languages/systems by using an LLM both to synthesize a distilled input prompt from documentation and examples and to generate and mutate inputs from it. CodaMOSA [12] invokes an LLM only after search-based test generation’s coverage plateaus, asking it for example test cases to escape the plateau – i.e., the LLM is invoked because coverage feedback is available and has stalled. This paper’s setting differs along the same axis as its predecessor [1]: no coverage signal is available at any point, and the LLM’s job is to revise a reusable, executable generator against feedback a black-box harness can actually produce, not to emit one-off inputs.

e) Differential testing.: Running multiple independent implementations of the same specification against identical inputs and comparing their outputs, rather than checking each against a hand-written oracle, dates to McKeeman’s differential testing [13]. This idea is what makes RQ2’s oracle possible at all in a setting with no crash to detect: four independently-implemented Dart JSON deserialization paths sharing one formal schema give a comparison target for free, the same way McKeeman compared independent compiler implementations against each other rather than against a formal semantics. Where classic differential testing compares whole programs given the same input, our four paths share almost their entire input pipeline (Section III’s shared JSON-syntax gate) and differ only in schema-level decoding, which is a narrower and, we find, more failure-prone surface than differential testing is usually applied to.

f) Positioning.: This paper is a direct empirical follow-up to [1], which evaluated the identical refinement-loop mechanism on two independently implemented C JSON parsers under a sanitizer-based crash oracle and named a memory-safe target as future work along a

materially different generalization axis. We take up exactly that axis: same loop, same sandboxing discipline, same reproducibility tooling where none of it is C- or format-specific, retargeted at a language where the entire premise of a crash oracle does not hold, using differential divergence [13] and ground-truth numeric round-tripping as replacements. Unlike TitanFuzz, Fuzz4All, and CodaMOSA, the LLM here never sees coverage data, execution traces, or usage examples – only a schema description and the same three-signal proxy discipline (here, a divergence score in place of acceptance rate) a human maintaining a Hypothesis strategy by hand would have to reason about manually.

III. Methodology

A. Research questions

RQ1: reusing the original acceptance-rate proxy and refinement loop unmodified in mechanism, does LLM-guided refinement raise the fraction of generated documents a Dart JSON decoder accepts as syntactically valid, relative to a static grammar-only baseline, the way it did for a native C target? RQ2: across four independently-implemented Dart JSON deserialization paths, does malformed or boundary-case input surface behavioral divergence, and does refinement – once its prompt is corrected to make divergence the explicit objective – find more of it than a static, schema-aware generator? RQ3: does numeric decoding across these paths preserve exact integer values at magnitudes JSON permits syntactically but a 64-bit integer type does not, and how does this generalize once measured at scale?

B. Target set

Four Dart [14] JSON deserialization paths over one shared schema (Section III-C), chosen for being how a real Flutter codebase actually deserializes JSON today, not for novelty: (A) hand-written parsing via `dart:convert`’s `jsonDecode` plus explicit `as` casts, representing code with no generated safety net; (B) `json_serializable` 6.14.1 [15], the most widely used Dart codegen serialization package; (C) `freezed` 4.0.1 [16], configured to delegate JSON leaf-decoding to `json_serializable` so that any B-vs-C divergence is attributable to `freezed`’s own contribution (immutability, union discrimination) rather than a second independent JSON decoder; (D) `built_value` 8.13.0 [17], an independently-designed, typed-serializer-based model rather than an annotation-and-Map-based one like B/C, and consequently the path most likely to diverge for genuinely independent-implementation reasons. All four resolve together under one Dart 3.13.2 toolchain with no dependency conflicts; exact versions are pinned in a committed lockfile.

C. Shared schema

One logical record, hand-implemented once per path: `id` (a bare JSON integer, RQ3’s target field), `amount` (a

JSON string, carrying text byte-for-byte with no numeric coercion by any path, so it cannot exhibit RQ3’s failure mode – it exists purely to widen RQ2’s perturbation surface), `name` (nullable string), `status` (one of three enumerated string values, chosen to probe unrecognized-enum-value handling), `tags` (a list of strings), and `child` (one bounded level of recursive nesting of the same record).

D. Two coverage-free, crash-free oracles

Every generated document passes through a two-stage harness. Stage one, a shared structural gate: `base64`-decode the candidate bytes, decode as UTF-8, then attempt `jsonDecode`. Any failure here – invalid UTF-8, invalid JSON syntax, or a syntactically valid but non-object top-level value (a bare array, number, or null, which JSON’s grammar permits but no path can schema-decode) – is recorded once, uniformly, since all four paths share this exact entry point and would fail identically before ever diverging. This has a direct consequence for generator design (Section III-G): a byte-level, grammar-only mutator can vary only what this shared gate tests, so it cannot find divergence by construction – divergence can only exist at the schema stage. Stage two, four independent schema attempts, reached only if stage one succeeds: each of the four paths attempts to decode the same parsed object, independently caught so that one path’s exception can never prevent the other three from running on the same input. Each result is either accepted, with a canonical value (Section III-E), or rejected, with the concrete runtime type of the exception and whether that type is private – Dart’s leading-underscore naming convention marking a class as internal to its defining library, meaning calling code cannot name it in a catch clause and must catch a generic supertype instead. This operationalizes “undocumented failure surface” mechanically rather than via a hand-maintained list of each package’s documented exceptions.

From these four records, two divergence-based oracles are computed automatically, with no crash required: cross-path divergence (RQ2) is accept/reject when acceptance status differs across paths, or value when two or more accepting paths decode to different canonical values. Numeric round-trip divergence (RQ3) compares each accepting path’s canonical id against the exact ground-truth integer parsed from the candidate’s own source text – Python’s arbitrary-precision `int` serializes exactly via `json.dumps` with no float detour, so ground truth requires no special-cased generation, only ordinary integer drawing at a biased range. A genuine process-level fault (the harness killed by a signal, or a batch timeout) is this pipeline’s closest analogue to the original’s sanitizer-based crash tier, checked around a whole campaign rather than around one input (Section III-F).

E. Canonical value comparison

Comparing decoded values safely requires that comparison itself not reintroduce the exact class of error the oracle exists to detect. Each path’s typed result is converted, once, to a shared representation using only Dart primitives – strings, exact-precision integers, booleans, null, and nested maps/lists – and compared structurally, never as a floating-point value. Any reported divergence is therefore real by construction, not an artifact of the comparison logic.

F. Harness architecture

The original pipeline spawns one process per input, which is cheap for a single native binary but would be four times as expensive here and provides no fault-isolation benefit a memory-safe language cannot already give in-process (an exception is caught, not a process-ending fault, unless truly fatal). The harness is instead one AOT-compiled, long-lived process per campaign, reading newline-delimited JSON requests from stdin – one line per candidate, base64-encoded since a candidate may be invalid UTF-8 by design – and writing one newline-delimited JSON result per line, flushed immediately so a driver can still attribute a later process-level fault to the correct input despite one process now serving an entire campaign rather than one input.

G. Generators

RQ1 reuses the original pipeline’s grammar-derived generator (the ANTLR grammars-v4 JSON grammar [18], [19] translated by hand into Hypothesis combinators) completely unmodified, targeting the shared structural gate directly: “accepted” is defined exactly as the original grammar defines it – any value type at the top level, not only objects – specifically so this research question tests portability of the refinement mechanism on a like-for-like criterion, not a criterion this paper’s own schema happens to impose. RQ2/RQ3 require a schema-aware generator instead, since Section III-D’s shared gate makes a grammar-only mutator structurally incapable of finding schema-level divergence: a base valid-record generator draws each field at its correct type, with id biased toward magnitudes near double’s exact-integer boundary and 64-bit integer boundaries; an independently-probable perturbation layer then drops a field, substitutes a wrong-but-JSON-valid type, nulls a field regardless of nullability, substitutes an out-of-enum string, injects an unknown key, or forces deeper recursion than the base generator’s normal cap. For the LLM-refined arm, the identical @st.composite def generated_json(draw) -> bytes contract is used, with a schema description substituted for the grammar text in the prompt and, for RQ2, a divergence score (Section III-I) in place of acceptance rate as the quantity the prompt asks the model to maximize.

H. What transfers unmodified from the original pipeline

The original project’s AST-based sandbox over LLM-authored proposals (import restricted to hypothesis.strategies, a blocklist over named code-execution/I-O entry points, and a behavioral sanity check that the proposal actually produces bytes); its bounded iteration budget and error-recovery discipline (a failing proposal consumes an iteration slot and its exact error text is fed into the next prompt, rather than aborting the run); its Hypothesis reproducibility shim (pinned to the exact Hypothesis version whose private internal it patches); and its per-iteration artifact discipline (every prompt, proposal, and result persisted, nothing discarded) are all reused with no logic changes, only the target-specific campaign runner beneath them replaced. The structural-fingerprint diversity proxy is likewise reused verbatim: it is JSON-specific, not C-specific, and the format has not changed.

I. RQ2’s refinement prompt: a confirmed fix, not a guess

An initial prompt presented the divergence count alongside per-path rejection counts in one undifferentiated metrics object, with no relative weighting. Across three full seeded runs this produced zero divergence despite driving rejection counts up sharply – confirmed as systematic, not run-to-run noise, before any prompt change was made (Section IV-B). The prompt was revised to state an explicit numeric score (divergence count over total), to say plainly that a high rejection count alone is not the goal, and to add a methodological hint – vary one or two things about an otherwise-valid document rather than maximizing how broadly malformed a document looks – deliberately worded to avoid naming the specific fields or outcomes already known by hand, so that any resulting discovery would be the loop’s own. The fix was validated on a single held-out seed before being adopted for the reported comparison (Section IV-B).

J. Follow-up arms on the RQ2 gap

Five further refined arms, each on fresh seeds disjoint from every other run and otherwise identical to the main refined arm, test why refinement underperforms on RQ2. JSON allowed (seeds 100–104): the sandbox forbids importing json, forcing the LLM to hand-build JSON text; this arm permits import json (via a separate sandbox module, so the original validator is never modified), isolating generation overhead from targeting. Category feedback (seeds 200–204): in addition to the score, the prompt reports how many divergence-causing documents exhibited each perturbation category – missing field, null override, wrong type, unrecognized enum, boundary-magnitude id, extra key, deep nesting – classified automatically from each document’s shape and never naming specific fields, preserving Section III-I’s discipline against planting the answer. Mutate source (seeds 400–404): every other arm discards the previous iteration’s generator; this one shows

the proposer its own previous source and asks it to revise it incrementally. Because a held-out seed showed a buggy revision persisting for the rest of a run, it carries a revert-on-regression guard: if the syntactically valid fraction drops more than 10 percentage points below the run’s best so far, the next iteration mutates from the best-known source instead, with a note telling the proposer what regressed.

Knowledge parity (seeds 500–509): the static generator’s perturbation catalog was designed after the thirteen-case characterization of the four paths, while every refined arm above sees only the schema and its score. This arm adds to the score-only prompt exactly what that designer knew beforehand, as four plain observations with no generator code: invalid JSON is rejected identically by all four paths; built_value accepts a missing tags while the others reject it; wrong types, nulls, and unknown enum values are rejected by all four, while extra keys and missing nullable fields are accepted by all four; and json_serializable/freeze decode id via (json['id'] as num).toInt(), which saturates out-of-range values. Stronger proposer (seeds 605–609): the score-only arm with gpt-4.1 in place of gpt-4.1-mini – same family, same API parameters, not a reasoning model – testing whether a more capable model discovers that knowledge from the score alone. Because gpt-4.1 routinely appends prose after its fenced code block, which the original response parser does not strip (a first run on seeds 600–604 lost 11 of 25 iterations to this), this arm extracts the first fenced block; the fix is opt-in, affected at most 2 iterations of any gpt-4.1-mini arm, and favors gpt-4.1. For both arms, seeds, tests, and decision rules were fixed in the released design document before any run.

K. Experimental design

All three research questions use seeded, independent runs (Hypothesis’s draw stream reseeded once per run via the reproducibility shim above), with $n=5$ per arm under a fixed API budget, reported as such throughout – the same discipline the original project applied to its own smaller-sample ablation. RQ1 compares a static baseline (one iteration, the grammar generator alone) against five refined runs (five iterations of up to 500 examples each, gpt-4.1-mini, temperature 0.2). RQ2 compares a static baseline (the schema-aware generator, one campaign of 500 examples per run) against five refined runs under the corrected prompt. RQ3 requires no separate campaign: every schema-evaluated record collected for RQ2 already carries each path’s canonical id and the document’s own source text, so RQ3’s ground-truth comparison is computed as an analysis pass over RQ1/RQ2’s existing data rather than a third parallel generator/harness/campaign pipeline.

TABLE I
RQ1: acceptance rate, $n=5$ seeded runs per arm.

Metric	Value
Baseline, per-run (%)	50.0, 50.0, 50.0, 50.0, 50.0
Refined, per-run (%)	68.6, 97.3, 98.0, 98.1, 98.2
Baseline mean (stdev)	50.0 (0.0)
Refined mean (stdev)	92.0 (13.1)

IV. Evaluation

All reported p -values are exact, two-sided Mann-Whitney U tests (scipy.stats.mannwhitneyu, method="exact"). At $n=5$ per arm, $C(10, 5) = 252$ and exactly two partitions produce complete rank separation, so $p = 2/252 \approx 0.00794$ in every complete-separation case below – the smallest p this design can produce, not a coincidence between unrelated effects.

A. RQ1: does the refinement mechanism transfer?

Table I reports acceptance rate – the fraction of generated documents that are syntactically valid JSON by the same criterion the original grammar uses, any value type at the top level – for five seeded baseline runs and five seeded refined runs. The baseline is exactly 50.0% on every run by construction (its input stream deterministically alternates a grammar-valid generator with a deliberately corrupted variant of it). Every refined run exceeds every baseline run: complete rank separation, exact $p \approx 0.00794$. One refined run’s lower rate (68.6%) traces to a seed on which only one of five iterations produced usable data, the rest consumed by ordinary LLM code-generation failures (e.g. one proposal built an unbounded-depth strategy that hit Python’s recursion limit) – logged and handled exactly as the original pipeline’s error-recovery design specifies.

This reproduces the original study’s cJSON result in shape – 58.2% to 97.1% there, 50.0% to a 92.0% mean here – on a different runtime and language, with the identical refinement-loop mechanism. We read this as confirmation that the mechanism, not merely the original numeric result, generalizes.

B. RQ2: does refinement find more divergence than a static generator?

The first refinement prompt (Section III-I) was run across three full seeded runs (seeds 0–2) before any revision: zero cross-path divergence in every resulting iteration-campaign, despite per-path rejection counts climbing sharply (up to 208 of 264 schema-evaluated documents on one run). The revised prompt was validated on one held-out seed (11.8% divergence in its best iteration) before being adopted for the five-seed comparison in Table II.

Every baseline run exceeds every refined run: complete rank separation, exact $p \approx 0.00794$, in the opposite direction from RQ1. Part of the gap is generation overhead – the sandbox forbids importing a JSON library, so

TABLE II

RQ2: cross-path divergence rate (% of all generated documents), seeded runs per arm ($n=5$ unless noted). p : exact two-sided Mann-Whitney U against the score-only refined arm. Every arm except parity is in complete separation below the static baseline ($p \approx 0.0079$). *Knowledge parity, $n=10$.

Arm	Per-run (%)	Mean (sd)	p
Static baseline	19.8, 21.8, 22.2, 23.4, 23.8	22.2 (1.57)	0.0079
Score-only	1.3, 1.45, 2.75, 3.6, 5.53	2.93 (1.74)	–
JSON allowed	0.05, 2.2, 2.36, 2.36, 4.05	2.20 (1.42)	0.69
Categories	0.0, 0.0, 1.0, 1.2, 2.55	0.95 (1.05)	0.032
Mutate source	0.0, 0.30, 0.65, 2.93, 4.75	1.73 (2.05)	0.31
gpt-4.1	0.0, 0.36, 0.60, 2.52, 6.05	1.91 (2.51)	0.31
Parity*	0.0, 17.6, 22.1, 26.4, 27.7, 31.0, 36.6, 38.7, 41.6, 45.9	28.8 (13.45)	0.013

the LLM hand-builds JSON text, and the refined arm’s schema-evaluated fraction ranged 41.9%–83.7% across runs versus the baseline’s fixed 100% – but restricting to schema-evaluated documents still leaves 4.10% against 22.2%. Pooling every schema-evaluated record across the five refined seeds, the loop found exactly the same two divergence patterns the static generator found – the two defects in Section IV-G – never a qualitatively new one, at a lower rate throughout.

C. Why does refinement underperform? Three follow-ups

The lower rows of Table II report the three follow-up arms of Section III-J.

Generation overhead or targeting? With import json permitted, every one of the 24 iterations that produced data reached 100% (in one case, affected by an unrelated encoding issue, 82.5%) schema-evaluated output – the overhead is essentially eliminated – yet the divergence rate did not improve: 2.20% mean, statistically indistinguishable from the constrained sandbox ($p \approx 0.69$). Removing the overhead entirely closes no measurable part of the gap; the shortfall is a targeting problem.

Does richer feedback help? Category feedback did measurably worse: 0.95% mean against the score-only prompt’s 2.93%, both over all documents and restricted to schema-evaluated ones ($p \approx 0.032$ in each case, ruling out proposal-failure noise). Two of the five seeds found zero divergence in their only successfully validated iteration. We have no confirmed account of why; a plausible, untested reading is that a longer prompt enumerating seven categories invites more elaborate, more bug-prone generators that try to target each category explicitly.

Does mutating, instead of rewriting, help? A first held-out seed (300) without the guard showed the failure mode the guard addresses: a revision in iteration 3 broke JSON syntax for about 27–30% of output and persisted to the end of the run (0.76% divergence). With the guard, the same seed reached 5.33% overall, and 10.91% of schema-evaluated documents – above every earlier refined run on that metric – which justified the $n=5$ comparison. It did not replicate: 1.73% mean, indistinguishable from both

score-only ($p \approx 0.31$) and category feedback ($p \approx 0.69$). The held-out seed cleared one early proposal failure quickly and then had clean iterations to compound on; seed 400, by contrast, failed on four of five iterations (the proposer returning an example document instead of a strategy) and found 0%. Mutation’s capacity to compound depends on reaching a working generator early, and a run that does not is left with little signal regardless of the guard.

Three independent interventions – one on generation, one on feedback, one on loop structure – all land in the same regime: statistically indistinguishable from each other or worse, and an order of magnitude below the static baseline.

D. Knowledge or search? Parity and a stronger proposer

The last two rows of Table II separate the two explanations the follow-ups leave open: whether the loop lacks the ability to search for divergence, or lacks the domain knowledge the static generator’s designer had. Knowledge parity. Given the four observations of Section III-J and nothing else new, the loop’s divergence rate rises by an order of magnitude. A first five seeds were suggestive but not significant against score-only ($p \approx 0.15$; one run lost four of five iterations to proposal failures and scored 0%), so we ran a replication on five fresh seeds under a decision rule fixed in advance: all five replication runs (22.1–41.6%) exceed every score-only run ($p \approx 0.0079$), and the pre-declared primary test on all ten seeds gives $p \approx 0.013$ (Cliff’s $\delta = 0.80$). With parity of knowledge, the loop is statistically indistinguishable from the static baseline over all documents (28.8% vs. 22.2% mean, $p \approx 0.16$) and significantly above it over schema-evaluated documents (40.7% vs. 22.2%, $p \approx 0.013$). It exploits what it is told but finds nothing beyond it: every divergence it produced is one of the same two patterns. Stronger proposer. Without that knowledge, gpt-4.1 is no better than gpt-4.1-mini (1.91% vs. 2.93% mean, $p \approx 0.31$) and remains far below the baseline ($p \approx 0.0079$). The RQ2 gap is therefore a knowledge gap – what a thirteen-case manual characterization supplies – that neither five iterations of score feedback nor a larger model closes, and that the loop closes fully once given it.

E. RQ3: does numeric decoding preserve exact values?

Table III reports the ground-truth id comparison (Section III-D) computed as an analysis pass over schema-evaluated records already collected across the RQ2 campaigns (static baseline, score-only refined arm, and JSON-allowed arm, pooled) – no separate generator, harness, or campaign was needed. Because the schema’s recursive child field makes each accepting path’s canonical value mirror the full decoded tree, every nesting depth is another independent check; walking all of them gives $N=43,336$.

Manual parsing and built_value never silently mismatch: across ten thousand accepted decodes each,

TABLE III

RQ3: exact-integer round-trip on the id field, $N=43,336$ checks (all nesting depths, pooled across RQ2 campaigns).

Path	Accepted	Mismatched	Rate
Manual	10,070	0	0.0%
json_serializable	11,449	649	5.7%
frozen	11,449	649	5.7%
built_value	10,368	0	0.0%

they either decode id exactly or reject the input outright. json_serializable and frozen are identical to each other (frozen delegates this field’s decoding to json_serializable’s generated code) and both show a 5.7% silent mismatch rate. Every observed mismatch saturates to exactly int64::MAX or int64::MIN – e.g. an input id of $-9,223,372,036,854,775,810$ decodes to $-9,223,372,036,854,775,808$ – because the generated code calls `(json['id'] as num).toInt()` rather than `json['id'] as int`: an out-of-range integer literal is first promoted to a double by `dart:convert`, and `toInt()` on an out-of-range double saturates rather than throwing.

The rate is not uniform across nesting depth: 3.4% at depth 0 ($n=8,198$), 7.7% at depth 1 ($n=2,377$), 20.7% at depth 2 ($n=590$), and 21.2% at depth 3 ($n=240$), with too few samples beyond ($n \leq 23$ per level) to read further. We have no confirmed mechanism; a plausible, untested explanation is that the generator’s perturbation and recursion machinery are not independent, so deeply nested documents disproportionately carry boundary-magnitude id values. We report the pattern descriptively.

F. A side study: does frozen ever diverge from json_serializable?

In every result above, frozen and json_serializable are identical, because the Record schema never exercises frozen’s own native discriminated-union dispatch (Section III-B). A small, deliberately qualitative side study – not part of the RQ1–RQ3 statistics – implemented a two-variant union, `Payload.text(String) / Payload.number(num)` discriminated by a “type” key, once per path in each framework’s idiomatic style (frozen: `@Frozen(unionKey: 'type')`; the others: hand-written dispatch, since neither json_serializable nor built_value has native polymorphism), and ran thirteen hand-picked documents through all four paths. frozen diverges from json_serializable for the first time, but at the exception-type level: every invalid discriminant (unknown, missing, null, wrong-type, differently-cased) is rejected by all four paths, yet frozen raises the public, field-attributed `CheckedFromJsonException` while the other three raise a generic `ArgumentError`. Invalid leaf values show no B-vs-C difference, and no accept/reject or value divergence appeared in this sample. Practically, json_serializable’s manual dispatch gives callers a less specific exception than frozen’s native unions for the failure mode polymorphic

JSON is most likely to hit; running this schema through the automated pipeline at scale is future work.

G. Two concrete, practitioner-actionable defects

Both were found by the static generator alone, with no LLM involved. (1) Silent integer saturation (RQ3): a Flutter application using json_serializable or frozen to decode an identifier that can exceed 64-bit range (a snowflake ID, an external database key) silently receives `int64::MAX` or `int64::MIN` rather than an exception; the same `toInt()` call also silently truncates a fractional literal (1.9 decodes to 1), which we confirmed by hand. (2) Silent default on a missing collection: built_value decodes a document whose required tags list is absent (or explicitly null) as an empty list, while the other three paths reject it. Missing scalar fields (id, amount, status) are rejected by every path including built_value, so the finding is scoped to its auto-created collection builders, not general laxness. Both defects reproduce with the released harness, and we have reported both to the respective maintainers.

V. Discussion

A. Interpreting the RQ1/RQ2 contrast

The central empirical fact of this paper is not either number in Section IV alone, but that they point in opposite directions under an otherwise identical mechanism: the same loop, proposer model, iteration budget, and sandboxing discipline raise acceptance rate on a like-for-like replication of the original objective (RQ1) and underperform a simple static generator on a differential objective (RQ2). “LLM-guided refinement helps” is not a property of the pipeline in general, but of the pipeline given a specific objective. Acceptance rate against a single decoder is a smooth signal: almost any change that makes a document more JSON-shaped moves it in the right direction, and an LLM has seen enormous amounts of well-formed JSON. Cross-implementation divergence requires a document that is simultaneously valid enough to clear one path and malformed in one specific, narrow way that trips another – a needle-in-haystack search over a far smaller region of the input space.

The follow-ups and the parity arm together locate where the difficulty lies. Removing generation overhead changed nothing, so the problem is not producing valid JSON. Richer feedback and incremental mutation did not help, and a larger model did no better. But handing the loop four observations from a thirteen-case manual characterization raised its rate by an order of magnitude, to the static generator’s level and above it on valid documents (Section IV-D). The RQ2 gap is therefore not an inability to search or to implement a generator: it is missing domain knowledge, which score feedback over five iterations does not recover – plausibly, though untested, because a proposer that scores near zero sees too few divergent documents to generalize from. Once told, the loop exploits the knowledge fully, but discovers nothing

beyond it – every divergence in every arm is one of the same two patterns.

For practitioners, this reframes the cost comparison. The static generator did not beat the loop because hand-written generators are inherently better; it won because it encoded what a thirteen-case manual characterization revealed. Feeding that characterization to the loop recovers the same performance with no hand-written generator at all, which suggests a practical workflow: characterize a handful of cases by hand, then let refinement exploit the findings – rather than expecting refinement to replace the characterization.

B. Reporting a corrected prompt, not just a final one

The initial RQ2 prompt drove up an unrelated metric (per-path rejection counts) while finding zero divergence across three full seeded runs – a proposer optimizing “move the numbers shown to me” when the prompt does not identify which number is the objective. This mirrors, on a different metric, the original study’s concern about acceptance-rate over-optimization crowding out boundary-probing behavior. The failure mode is informative on its own, and silently iterating past it would have understated how easily an LLM-guided loop optimizes the wrong thing.

C. Threats to validity

Construct validity. The divergence oracle detects disagreement between the four instrumented paths, not whether any path is “correct” relative to a specification. RQ3’s ground-truth oracle is scoped to the one field (id) that is a bare JSON number; the string-typed amount cannot exhibit numeric precision loss.

Internal validity. The corrected RQ2 prompt was validated on a held-out seed before the five-seed comparison, to avoid tuning on the reported data, but we tried no further prompt variants beyond it and the follow-ups, so a different, equally principled design might close more of the gap. The knowledge-parity arm’s four observations were checked line by line against the characterization output and stated only what was known before the static generator was written, but any such paraphrase involves judgment about what to include. Its first five seeds were not significant on their own; the reported result rests on a replication whose seeds, primary test, and decision rule were fixed in advance, and we report both. Two of our pre-declared decision rules turned out not to cover every outcome (one read a non-significant result as evidence of no effect; one assumed a stronger model could only do better); the released design document records both gaps alongside the results.

Statistical validity. $n=5$ per arm is a small sample, chosen under a fixed API budget and reported as such. Complete separation at $n=5$ gives $p \approx 0.00794$, the smallest attainable value, not an effect size; we report

per-run rates in Tables I and II for that reason. Mutate-source in particular has high per-run variance (stdev 2.05, on par with its mean).

External validity. All results are on JSON, four pinned-version Dart deserialization paths, and two proposer models from one family (gpt-4.1-mini, and gpt-4.1 for one arm) at one temperature. We make no claim that the magnitudes of the RQ1, RQ2, or follow-up effects generalize to other formats, packages, versions, or model families, including reasoning models.

D. When this approach is, and is not, the right tool

The divergence oracle is informative only when multiple independent implementations of the same specification are in real concurrent use, as with Dart’s several competing codegen frameworks. A target with one canonical implementation has no comparison partner; the numeric round-trip oracle or an exception-contract oracle (the private-exception classification of Section III-D) would apply instead.

E. Future work

The parity result raises a sharper question than the one this paper started with: can the loop acquire the knowledge itself, for instance by running a small, automated characterization step before refinement, or by reading the generated deserialization code the way the human designer did? Whether a reasoning model or a different model family discovers it unaided is also open. A larger- n replication would tighten every comparison beyond what $n=5$ allows. Finally, running the union schema (Section IV-F) through the automated oracle and seeded campaigns would establish whether its exception-type difference holds at scale rather than on thirteen cases.

VI. Conclusion

We retargeted an LLM-guided, coverage-free grammar-fuzzing pipeline at a memory-safe language, where its sanitizer-based crash oracle has no meaning, replacing that oracle with cross-implementation differential divergence across four Dart JSON deserialization paths and a ground-truth numeric round-trip check. The unchanged refinement loop reproduces the original acceptance-rate result on this new target (50.0% to a 92.0% mean), but underperforms a static, schema-aware generator by an order of magnitude on the differential objective (2.9% vs. 22.2%). Generation overhead, feedback richness, loop structure, and model size do not explain the gap; domain knowledge does. Given four observations from the manual characterization the static generator was designed from, refinement matches that generator, yet never finds a divergence it was not told about. For anyone putting an LLM in a fuzzing loop, the practical lesson is that refinement exploits domain knowledge well but does not

discover it: a small manual characterization is the high-leverage input, and refinement is the tool for exploiting it rather than a replacement for it.

Independently of any LLM, the static generator’s campaigns expose two defects directly actionable for Flutter teams: `json_serializable` and `freezed` silently saturate out-of-range integers to `int64` bounds, and `built_value` silently defaults a missing required list to empty. We release the full pipeline, harness, and every campaign artifact underlying Tables I–III, so that each result can be checked, extended to a larger sample, or falsified directly.¹

References

- [1] Anonymous, “Agentic grammar fuzzing: LLM-guided grammar refinement for parser testing,” 2026, anonymized for double-blind review.
- [2] M. Zalewski, “American fuzzy lop,” <https://lcamtuf.coredump.cx/afl/>, 2014, accessed 2026.
- [3] A. Fioraldi, D. C. Maier, H. Eißfeldt, and M. Heuse, “AFL++: Combining incremental steps of fuzzing research,” in 14th USENIX Workshop on Offensive Technologies (WOOT 20). USENIX Association, 2020.
- [4] V. J. M. Manès, H. Han, C. Han, S. K. Cha, M. Egele, E. J. Schwartz, and M. Woo, “The art, science, and engineering of fuzzing: A survey,” *IEEE Transactions on Software Engineering*, vol. 47, no. 11, pp. 2312–2331, 2021.
- [5] P. Purdom, “A sentence generator for testing parsers,” *BIT Numerical Mathematics*, vol. 12, no. 3, pp. 366–375, 1972.
- [6] J. Wang, B. Chen, L. Wei, and Y. Liu, “Superion: Grammar-aware greybox fuzzing,” in *Proceedings of the 41st International Conference on Software Engineering (ICSE)*, 2019, pp. 724–735.
- [7] C. Aschermann, T. Frassetto, T. Holz, P. Jauernig, A.-R. Sadeghi, and D. Teuchert, “NAUTILUS: Fishing for deep bugs with grammars,” in *Proceedings of the 26th Network and Distributed System Security Symposium (NDSS)*, 2019.
- [8] K. Claessen and J. Hughes, “QuickCheck: A lightweight tool for random testing of Haskell programs,” in *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP)*, 2000, pp. 268–279.
- [9] D. R. MacIver, Z. Hatfield-Dodds, and Many other contributors, “Hypothesis: A new approach to property-based testing,” *Journal of Open Source Software*, vol. 4, no. 43, p. 1891, 2019.
- [10] Y. Deng, C. S. Xia, H. Peng, C. Yang, and L. Zhang, “Large language models are zero-shot fuzzers: Fuzzing deep-learning libraries via large language models,” in *Proceedings of the 32nd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA)*, 2023, pp. 423–435.
- [11] C. S. Xia, M. Paltenghi, J. L. Tian, M. Pradel, and L. Zhang, “Fuzz4All: Universal fuzzing with large language models,” in *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering (ICSE)*, 2024.
- [12] C. Lemieux, J. P. Inala, S. K. Lahiri, and S. Sen, “CodaMOSA: Escaping coverage plateaus in test generation with pre-trained large language models,” in *Proceedings of the 45th International Conference on Software Engineering (ICSE)*, 2023.
- [13] W. M. McKeeman, “Differential testing for software,” *Digital Technical Journal*, vol. 10, no. 1, pp. 100–107, 1998.
- [14] “The Dart programming language,” <https://dart.dev>, 2026, SDK version 3.13.2, accessed 2026.
- [15] “`json_serializable`: Automatically generate code for converting to and from JSON,” https://pub.dev/packages/json_serializable, 2026, version 6.14.1.
- [16] R. Rousset, “`freezed`: Code generation for immutable classes,” <https://pub.dev/packages/freezed>, 2026, version 4.0.1.
- [17] “`built_value`: Value types with builders, Dart classes as enums, and serialization,” https://pub.dev/packages/built_value, 2026, version 8.13.0.
- [18] “ANTLR grammars-v4: JSON grammar,” <https://github.com/antlr/grammars-v4/blob/master/json/JSON.g4>, 2024.
- [19] T. Parr, *The Definitive ANTLR 4 Reference*, 2nd ed. Pragmatic Bookshelf, 2013.

¹<https://anonymous.4open.science/r/agentic-fuzzing-dart-json-4656/>